

Dijkstra's Algorithm

Goal:

Given a graph where each edge has a non-negative weight, the algorithm computes the shortest (minimum total weight) path from a specified source vertex to every other vertex. It builds a “shortest path tree” where each vertex records the best way to reach it from the source.

Core Idea:

Dijkstra’s algorithm works by iteratively “relaxing” edges (i.e., updating the best known distances to vertices) and using a priority queue (min-heap) to always expand the vertex with the smallest tentative distance next.

Key Properties (Instance Variables)

- **distTo[] :**
 - *Purpose:* Holds the current shortest known distance from the source vertex to each vertex in the graph.
 - *Initialization:* Set to positive infinity (`Double.POSITIVE_INFINITY`) for all vertices except the source, which is set to 0.0.
- **edgeTo[] :**
 - *Purpose:* For every vertex, it stores the last edge used to reach that vertex on the current shortest path from the source.
 - *Use:* This array is used later to reconstruct the actual path from the source to any vertex.
- **pq (IndexMinPQ):**
 - *Purpose:* A priority queue that orders vertices by their current `distTo` value.
 - *Function:* Helps efficiently determine the next vertex to process (the one with the smallest tentative distance).

Methods and Their Purposes

1. Constructor: `DijkstraSP(EdgeWeightedDigraph G, int s)`

- **Purpose:**

Initializes the algorithm by computing the shortest paths from the source vertex `s` to every

other vertex in the graph G .

- **Key Steps:**

- **Input Validation:**

- Loops over all edges to ensure that none have a negative weight. If any negative edge is found, an exception is thrown.

- **Initialization:**

- Creates the `distTo` array with all values set to infinity except for the source (set to 0.0).
 - Sets up the `edgeTo` array to keep track of the path.

- **Priority Queue Setup:**

- Inserts the source vertex into the priority queue with its distance (0.0).

- **Main Loop (Relaxation Process):**

- Continues while the priority queue is not empty:
 - Removes the vertex v with the smallest `distTo` value.
 - Iterates through all edges leaving v and “relaxes” each edge using the `relax` method.

- **Optimality Check:**

- After computing the paths, the algorithm optionally checks that the computed paths satisfy the shortest path optimality conditions using the `check` method.

2. `relax(DirectedEdge e)`

- **Purpose:**

Examines an edge e from vertex v to vertex w to see if the known path to w can be improved by going from v through e .

- **Process:**

- It checks if `distTo[w] > distTo[v] + e.weight()`.
 - If true, it updates:
 - `distTo[w]` to the new, lower distance.
 - `edgeTo[w]` to the current edge e (this edge becomes part of the shortest path to w).
 - It then updates the priority queue:

- If w is already in the queue, its key (distance) is decreased.
- If not, w is inserted into the queue with its new distance.

3. `distTo(int v)`

- **Purpose:**

Returns the shortest distance from the source vertex to vertex v .

- **Usage:**

- After the algorithm has run, you can call this method to retrieve the minimum distance to any vertex.

4. `hasPathTo(int v)`

- **Purpose:**

Indicates whether there is any path from the source vertex to vertex v .

- **Logic:**

- It returns `true` if `distTo[v]` is less than infinity, meaning the vertex is reachable from the source.

5. `pathTo(int v)`

- **Purpose:**

Provides the actual shortest path (as a sequence of edges) from the source vertex to vertex v .

- **Process:**

- If no path exists (i.e., `hasPathTo(v)` is false), it returns `null`.
- Otherwise, it builds the path by:
 - Starting at vertex v and following `edgeTo` backwards until reaching the source.
 - Using a stack to reverse the order, so the path is returned from source to destination.

6. `check(EdgeWeightedDigraph G, int s)`

- **Purpose:**

A private method used to verify that the computed shortest path tree satisfies all optimality conditions. It helps to ensure the algorithm's correctness.

- **What It Checks:**

- All edge weights are non-negative.

- For the source, `distTo[s]` should be 0 and `edgeTo[s]` should be `null`.
- For every vertex, if there is an edge in `edgeTo`, then the distance to that vertex should be exactly equal to the distance to its predecessor plus the weight of the connecting edge.
- Every edge in the graph should not offer a shortcut that would reduce the computed `distTo` value (i.e., `distTo[w]` should be no more than `distTo[v] + weight` for each edge from `v` to `w`).

7. `validateVertex(int v)`

- **Purpose:**

A helper method that checks if a given vertex index `v` is valid (i.e., it lies between 0 and the number of vertices minus one).

- **Outcome:**

- If the vertex is not valid, it throws an `IllegalArgumentException`.

8. `main(String[] args)`

- **Purpose:**

Serves as a test client to run the algorithm from the command line.

- **Steps in `main`:**

- Reads a graph from a file (provided as the first command-line argument).
- Reads the source vertex (provided as the second command-line argument).
- Instantiates the `DijkstraSP` object to compute the shortest paths.
- Iterates over all vertices and prints:
 - The shortest distance from the source to each vertex.
 - The path (i.e., the sequence of edges) if one exists; otherwise, it prints that no path exists.

How to Translate This to Java or Python

For Java:

- Use similar class structures, arrays for distances and paths, and an index-based priority queue.
- Implement helper methods for relaxing edges and validating vertices.
- Follow object-oriented practices to encapsulate the graph, edge, and priority queue

functionalities.

For Python:

- You can use lists (or dictionaries) for `distTo` and `edgeTo` .
- The priority queue can be implemented using the `heapq` module.
- Define a class (e.g., `DijkstraSP`) that contains methods for initialization, edge relaxation, and retrieving paths.
- Ensure you handle cases where a vertex is unreachable by returning `None` or an equivalent.
- Since Python does not have built-in type enforcement like Java, be mindful of using exceptions or checks to validate input (like vertex indices).

Summary

- **Properties:**

- `distTo` : Stores the best-known distances from the source.
- `edgeTo` : Keeps track of the last edge on the path to each vertex.
- `pq` : A priority queue for efficiently selecting the next vertex to process.

- **Core Methods:**

- **Constructor:** Initializes distances, validates inputs, and performs the relaxation process using a priority queue.
- **`relax()`:** Updates distances and paths if a better route is found via an edge.
- **`distTo()`, `hasPathTo()`, `pathTo()`:** Provide access to the computed shortest distances and paths.
- **`check()` & `validateVertex()`:** Ensure the algorithm's correctness and input validity.