

8-Puzzle Solver Using A* Search

Project Description

Overview:

This project implements a solver for the 8-puzzle problem (and its generalizations to an $n \times n$ board) using the A* search algorithm. In the 8-puzzle, eight numbered tiles and one blank are arranged on a 3×3 grid. The goal is to slide the tiles—by moving a neighboring tile into the blank space—until the board is arranged in row-major order (with the blank in the last position). Note that not every board configuration is solvable. This implementation uses a *twin board* strategy (swapping two non-blank tiles) to detect unsolvability.





Figure 1: 8-Puzzle transitions from initial to goal

Key Components

1. Board Data Type

- **Representation:**

The board is modeled as an immutable object containing an $n \times n$ grid. Tiles are stored where 0 represents the blank square.



board

```
3
1 0 3
4 2 5
7 8 6
```

string representation

Figure 2: Board representation and string representation

- **Methods:**

- **Constructor:** Initializes the board from an $n \times n$ array (or list of lists in Python).
- **toString() / __str__:** Returns a string representation that shows the board size and layout.
- **Heuristics:**
 - * *Hamming distance:* Counts the number of tiles out of place.
 - * *Manhattan distance:* Sums the vertical and horizontal distances from each tile to its goal position.




Figure 3: Hamming and Manhattan distances example

- **isGoal():** Checks if the board matches the goal state.
- **neighbors():** Generates all valid boards reachable by sliding a tile into the blank space.




Figure 4: Board and neighbors

- **twin():** Creates a board by swapping any pair of non-blank tiles (used to detect unsolvable puzzles).

2. Solver Data Type

- **Search Node:**

Each search node stores:

- A board.
- The number of moves made so far.
- A link to the previous node.
- A priority computed as the sum of the moves and the board's Manhattan distance.

- **A* Search:**
Uses a priority queue to always expand the node with the lowest priority. When the goal board is reached, the solution path is reconstructed.
- **Dual Search for Unsolvability:**
Two searches are performed in lockstep—one on the initial board and one on its twin. If the twin reaches the goal, the original board is unsolvable.
- **Optimizations:**
 - Caching of the Manhattan (or Hamming) distances.
 - Pruning redundant search nodes by skipping neighbors that revert to the previous board.

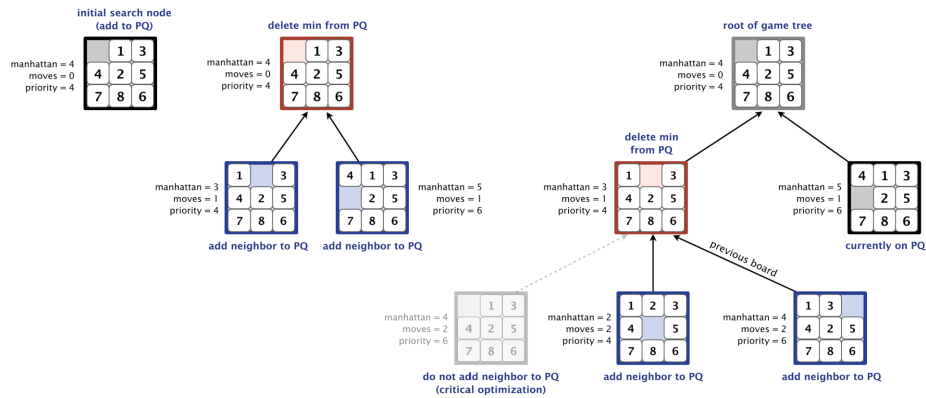


Figure 5: Search node expansions example

Hints on How to Solve the Problem

1. **Understanding the Board Representation:**
 - Model the board as an immutable data type so that once created, its state does not change. This helps in reliably comparing boards during the search.
 - Use a 2D array (or tuple of tuples) for representing the board configuration.
2. **Implementing Heuristics:**
 - **Hamming Distance:**
Count how many tiles (ignoring the blank) are not in their goal position.
 - **Manhattan Distance:**
For each tile, compute the distance from its current position to its target position (sum of horizontal and vertical differences).

- **Caching:**
Since the Manhattan distance is computed repeatedly, calculate it once for each search node and store it.
3. **Generating Neighbors:**
 - Identify the position of the blank tile.
 - Create new board configurations by sliding a neighboring tile into the blank space (considering up, down, left, and right moves).
 - **Optimization:** Do not enqueue a neighbor if it is the same as the previous board (to avoid cycling back).
 4. **A* Search Implementation:**
 - Use a priority queue (e.g., `heapq` in Python) to store search nodes.
 - Define each search node's priority as `moves + heuristic` (with the Manhattan heuristic preferred for better performance).
 - Expand the search by dequeuing the node with the smallest priority, then enqueue all its valid neighbors.
 5. **Detecting Unsolvable Boards:**
 - Generate a twin board by swapping two non-blank tiles.
 - Run A* search simultaneously on both the initial board and its twin. If the twin board reaches the goal state, then the initial board is unsolvable.
 6. **Reconstructing the Solution:**
 - When the goal state is reached in the search, backtrack using the linked search nodes to build the sequence of moves leading from the initial board to the goal.